

# Example of using Thread Class

```
public static void main(String[ ] args)
{
    System.out.println("Simple Thread Demonstration");
    System.out.println(" Extending Thread");
    for (int i = 0; i < 4; i++)
    {
        MyThread newThread = new MyThread( i );
        newThread.start( );
    }
}
```

Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

# Synchronization of Java Threads

- ❑ The previous examples shows independent, asynchronous threads (i.e., run at their own pace)
- ❑ Independent threads can run asynchronously if they only use local (private) data
- ❑ In many cases, threads i) share data, and/or ii) rely on the activities of other threads
- ❑ They must synchronize to share data
- ❑ Consider the following producer-consumer example

# Producer

```
public class Producer extends Thread {
    private CubbyHole cubbyhole;

    public Producer (CubbyHole c) {
        cubbyhole = c;
    }

    public void run ( ) {
        for (int i = 0; i < 10; i++) {
            cubbyhole.put(i);
            try {
                sleep ((int) (Math.random( ) * 100));
            } catch (InterruptedException e) {}
        }
    }
}
```

- ❑ Generates  $0 \leq i < 9$
- ❑ Stores it in a Cubbyhole
- ❑ Sleeps for a while
- ❑ Generates the next  $i$

```
public class CubbyHole {
    private int contents;

    public int get (int number)
    { ... }

    public void put (int number, int value)
    { ... }
}
```

# Consumer

```
public class Consumer extends Thread
{
    private CubbyHole cubbyhole;
    public Consumer (CubbyHole c) {
        cubbyhole = c;
    }
    public void run ( ) {
        int value = 0;
        for (int i = 0; i < 10; i++) {
            value = cubbyhole.get();
        }
    }
}
```

- ❑ Consumes the integers from the Cubbyhole Object
- ❑ As quickly as they become available

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Producer / Consumer Example

- ❑ CubbyHole is the shared data
- ❑ Ideally Consumer will get each value only once produced by Producer
- ❑ If Producer runs faster than Consumer,

Consumer #1 got: 3  
Producer #1 put: 4  
Producer #1 put: 5  
Consumer #1 got: 5

- ❑ Problem arises if Producer is quicker than Consumer
- ❑ Producer generates two numbers before Consumer consumes the first
- ❑ E.g. misses the number 4

# Producer / Consumer Example

Producer #1 put: 4  
Consumer #1 got: 4  
Consumer #1 got: 4  
Producer #1 put: 5

- ❑ If Producer is slower than Consumer we also have a problem
- ❑ Consumer gets two same numbers before Producer generates a new number
- ❑ E.g. gets the number 4 twice

Assignment Project Exam Help

<https://powcoder.com>

- ❑ Example of a race condition – reading and writing shared data concurrently; success depends on timing
- ❑ Example of critical section - Need synchronization



# Synchronization of Java Threads

- ❑ The **put** and **get** in the CubbyHole code are the so called *critical section (monitor region)*
- ❑ These should be marked as **synchronized**

## Assignment Project Exam Help

```
public class CubbyHole {  
    private int contents;  
    private boolean available = false;  
  
    public synchronized int get () { ... }  
    public synchronized void put (int value) { ... }
```

<https://powcoder.com>

Add WeChat powcoder

- ❑ Running synchronized method locks class
- ❑ Unlocked when method terminates

# Synchronization of Java Threads

- So now we have mutual exclusion, but how do Producer and Consumer cooperate to address the race condition?
- Consumer needs to wait until Producer has put new data, and Producer needs to notify Consumer after the data is put– and vice versa

- put and get need to wait on and notify each other of their activities

```
public synchronized int get {
```

```
    while (available == false) {  
        try { // wait for producer to put value  
            wait ( );  
        } catch (InterruptedException e) { }
```

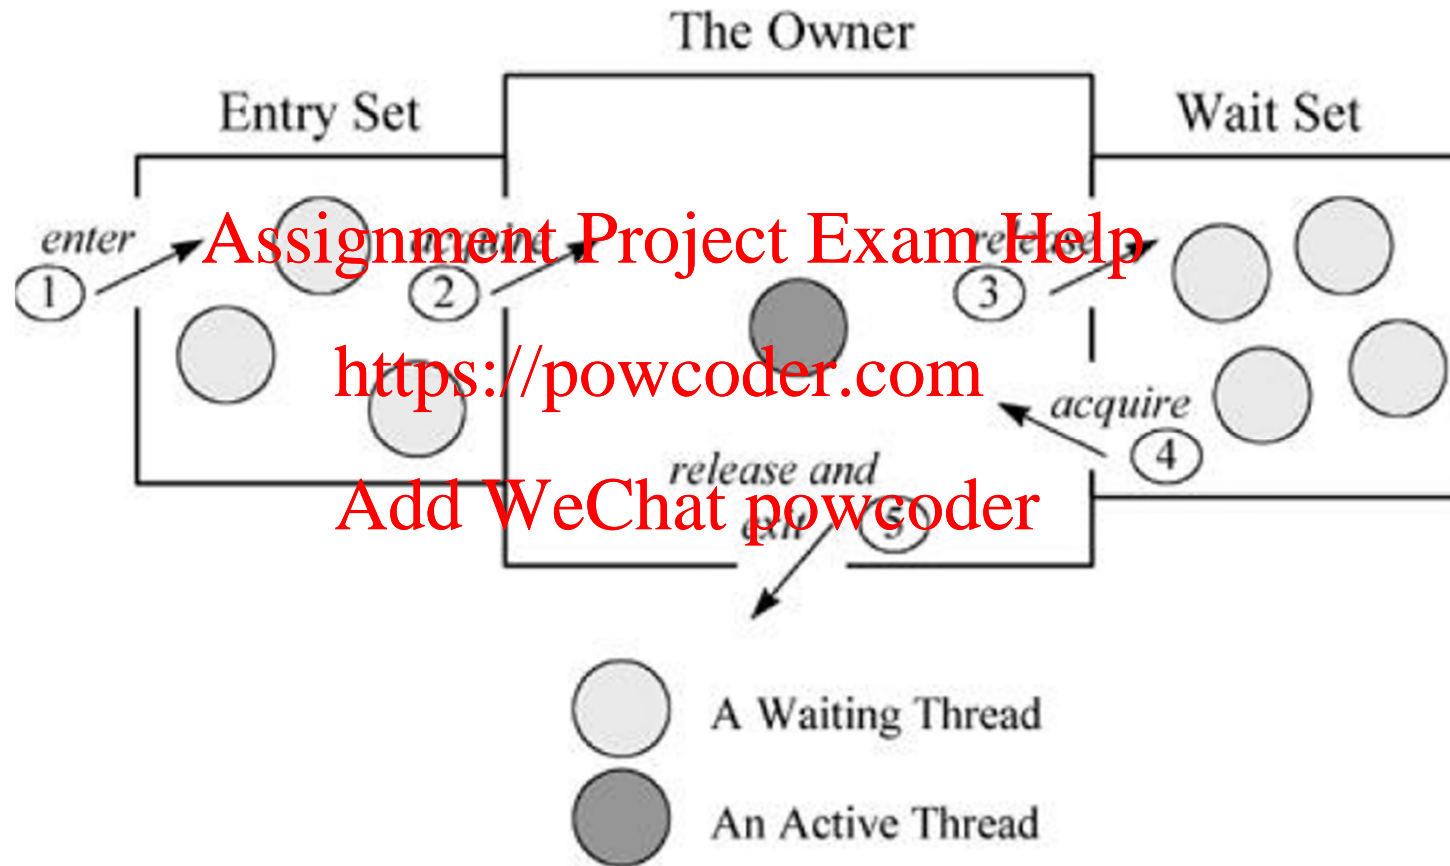
```
    }  
    available = false;  
    // Notify Producer value has been retrieved  
    notify ( );  
    return contents  
}
```

- available=false means the data has not been put

- wait releases lock



# Wait Releases Monitor



# Synchronization of Java Threads

- ❑ So now we have mutual exclusion, but how do Producer and Consumer cooperate?
- ❑ Consumer needs to wait until Producer has put something new, and Producer needs to notify Consumer after the data is put– and vice versa
- ❑ i.e put and get need to wait on and notify each other of their activities

```
public synchronized int get {  
  
    while (available == false) {  
        try { // wait for producer to put value  
            wait ( );  
        } catch (InterruptedException e) { }  
    }  
    available = false;  
    // Notify Producer value has been retrieved  
    notify ( );  
    return contents  
}
```

- ❑ available=false means the data has not been put

❑ wait releases lock

- ❑ Loops until the new data is available
- ❑ When consumer gets data, it does “notify” and Producer comes out of wait state and puts the next new data

# Synchronization of Java Threads

- ❑ So now we have locking, but how do Producer and Consumer cooperate?
- ❑ Consumer needs to wait until Producer has put something, at which point Producer needs to notify Consumer – and vice versa
- ❑ i.e put and get need to wait on and notify each other of their activities

```
public synchronized void put(int value){  
  
    while (available == true) {  
        try { // wait for consumer to get value  
            wait ( );  
        } catch (InterruptedException e) { }  
    }  
    contents = values;  
    available = true;  
    // Notify Consumer that value has been sent  
    notify ( );  
}
```

- ❑ available=true means the data has been put

- ❑ Loops until Producer ready with new value

- ❑ wait relinquishes lock

- ❑ When Producer put the data, it does notify and consumer comes out of wait state and get() can collect value

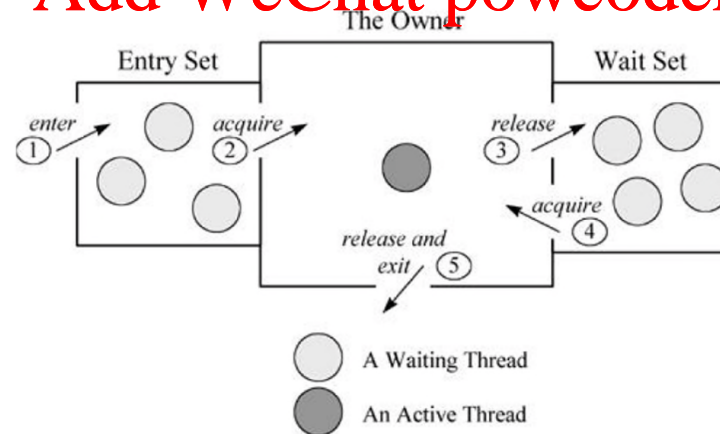
# Synchronization of Java Threads

- ❑ *notify* will allow one (rather than all) thread to be woken
- ❑ *notifyAll* wakes up all threads waiting on the monitor in question
- ❑ Awakened threads compete for lock (ownership of the monitor)
- ❑ Those threads that don't get the lock continue to wait

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Synchronization in Java (Monitor)

- Java provide **built-in** thread synchronization ability (through monitors), which is part of the programming language
- The compiler generate correct code to generate the monitor for each class/object and implement the monitor locks

For more information, see <https://powcoder.com>  
<http://java.sun.com/docs/books/tutorial/essential/threads>

Add WeChat powcoder



# Synchronisation in C

```
pthread_mutex_t my_mutex; /* declared in global area*/
```

```
void *Calc(void *);
```

```
main(){
```

```
    pthread_mutex_init(&my_mutex, NULL) /* before pthread_create*/
```

```
    for(ith=0; ith<NUM_THREADS; ith++) pthread_create( , Calc, );
```

```
}
```

```
void *Calc(void *param){
```

```
    pthread_mutex_lock(&my_mutex);
```

```
    critical section;
```

```
    pthread_mutex_unlock(&my_mutex);
```

```
}
```



# Synchronization in Java (Monitor)

- Java provide **built-in** thread synchronization ability (through monitors), which is part of the programming language
- The compiler generate correct code to generate the monitor for each class/object and implement the monitor locks
- Higher level and less error prone than forcing programmers to explicitly handle locks

For more information, see:  
<http://java.sun.com/docs/books/tutorial/essential/threads>

# High Performance Computing *Course Notes*

Assignment Project Exam Help

<https://powcoder.com>

## Message Passing Programming I

Add WeChat powcoder

Dr Ligang He



# Message Passing Programming

- Message Passing is the most widely used parallel programming model
- Message passing works by creating a number of processes, uniquely named, that interact by sending and receiving messages to and from one another (hence the message passing)
  - Generally, processes communicate through sending the data from the address space of one process to that of another
    - Communication of processes (via message, socket, pipe, files)
    - Communication of processes with a process (via global data area)
- Message passing programs can be based on standard sequential language programs (C/C++, Fortran), augmented with calls to library functions for sending and receiving messages

## Message Passing Interface (MPI)

- There are different **message passing models**; MPI is the most popular models (PVM is another one);
- MPI is a **specification**, not a particular implementation
  - specify a number of routines as well as the parameters that the routines should have
  - Does not specify how these routines are implemented, nor process startup, error codes, amount of system buffer, etc
- MPI **implementation** is a library, not a language
  - There are different MPI implementations: MPICH, LAM/MPI, OpenMPI
- In terms of the scope, Message passing model > MPI specification > MPI implementation

# OpenMP vs MPI

- MPI is used on distributed-memory systems
- OpenMP is used on shared-memory systems
- Both are explicit parallelism
- OpenMP is higher-level control
  - ❑ Compiler can automatically parallelise the codes when instructed
- MPI is lower-level control
  - ❑ Data partition, allocation and communication are conducted by programmers

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# A brief history of MPI

- ❑ Message-passing libraries developed for a number of early distributed memory computers
- ❑ By 1993 there were many vendor specific implementations
- ❑ By 1994 MPI-1 was proposed
  - ❑ Emphasize message passing
  - ❑ Has a static runtime environment
- ❑ By 1996 MPI-2 was finalized
  - ❑ includes new features such as
    - parallel I/O,
    - dynamic process management
    - remote memory operations
- ❑ 2012 MPI-3 is formalized



# The MPI programming model

- MPI standards -

  - MPI-1, MPI-2, MPI-3

- **Assignment Project Exam Help**  
Standard bindings - for C, C++ and Fortran.

- <https://powcoder.com>  
There are MPI bindings for Python, Java etc (non-standard).  
**Add WeChat powcoder**

- We will stick to the C/C++ binding, for the lectures and coursework.

# MPI functions

- MPI is a complex system comprising of numerous
- functions with various parameters and variants
- Six of them are indispensable, but can write a large number of useful programs already
- Other functions add other features, including flexibility (datatype), robustness (non-blocking send/receive), efficiency (ready-mode communication), modularity (communicators, groups) or convenience (collective operations, topology).
- In the lectures, we are going to cover most commonly encountered functions

# Intuitive Interfaces for sending and receiving messages

- Send(data, destination), Receive(data\_location, source)
  - minimal interface
- Not enough in some situations, we also need
  - Message matching – sender may send multiple messages to the same receiver, add message\_id at both send and receive interfaces
  - they become Send(data, destination, msg\_id), receive(data, source, msg\_id)
  - Message\_id
    - Is expressed using an integer, termed as *message tag*
    - Can differentiate the messages from the same process
    - Enable the messages to be processed in an ordered fashion

# How to express the data in the send/receive interfaces

- Early stages:
  - (address, length) for the send interface
  - (address, max\_length) for the receive interface
- They are not always good
  - the data may not occupy contiguous memory locations
  - Storing format for data may not be the same in heterogeneous platform
- Eventually, a triple (address, count, datatype) is used to express the data to be sent and (address, max\_count, datatype) for the data to be received
  - Reflecting the fact that a message contains much more structures than just a string of bits, For example, (vector\_A, 300, MPI\_REAL)
  - Programmers can construct their own datatype
- Now, the interfaces become send(address, count, datatype, destination, msg\_tag) and receive(address, max\_count, datatype, source, msg\_tag)

# How to distinguish messages

- Message tag is necessary, but not sufficient

- **Assignment Project Exam Help**  
So, communicator is introduced ...

<https://powcoder.com>

Add WeChat powcoder



# Communicators

- Messages are put into communication contexts
  - Contexts are allocated at run time by the MPI system
  - Each generated context is unique
- The processes in a MPI system are divided into groups
- The notions of context and process group are combined in a single object, which is called a *communicator*
  - A communicator consists of a group of processes and a communication context
  - The MPI library defines a initial communicator, MPI\_COMM\_WORLD, which contains all the processes running in the system
- Messages from different communicators can have the same tag
- So the send interface becomes send(address, count, datatype, destination, tag, comm)



# Status of the received messages

- The structure of the message status is added to the receive interface
- Status holds the information about source, tag and actual message size
  - In the C language, source can be retrieved by accessing `status.MPI_SOURCE`,
  - tag can be retrieved by `status.MPI_TAG` and
  - actual message size can be retrieved by calling the function `MPI_Get_count(&status, datatype, &count)`
- The receive interface becomes `receive(address, maxcount, datatype, source, tag, communicator, status)`

# How to express source and destination

- The processes in a communicator (group) are identified by ranks
- If a communicator contains  $n$  processes, process ranks are integers from 0 to  $n-1$
- Source and destination processes in the send/receive interface are the ranks

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## Some other issues

In the receive interface, tag can be a wildcard (MPI\_ANY\_TAG), which means any message will be received

**Assignment Project Exam Help**  
In the receive interface, source can also be a wildcard (MPI\_ANY\_SOURCE), which match any source  
<https://powcoder.com>

**Add WeChat powcoder**

# MPI basics

## First six functions (C bindings)

**MPI\_Send (buf, count, datatype, dest, tag, comm)**

*Send a message*

*buf* address of send buffer

*count* no. of elements to send ( $\geq 0$ )

*datatype* of elements

*dest* process id of destination

*tag* message tag

*comm* communicator (handle)

# MPI basics

## First six functions (C bindings)

**MPI\_Send** (buf, count, datatype, dest, tag, comm)

*Send a message*

*buf* address of send buffer

*count* no. of elements to send ( $\geq 0$ )

*datatype* of elements

*dest* process id of destination

*tag* message tag

*comm* communicator (handle)

# MPI basics

## First six functions (C bindings)

**MPI\_Send** (buf, count, datatype, dest, tag, comm)

*Send a message*

*buf* address of send buffer

*count* no. of elements to send ( $\geq 0$ )

*datatype* of elements

*dest* process id of destination

*tag* message tag

*comm* communicator (handle)



# MPI basics

## First six functions (C bindings)

**MPI\_Send**(buf, count, datatype, dest, tag, comm)

Assignment Project Exam Help

*Calculating the size of the data to be send ...*

<https://powcoder.com>

*buf* address of send buffer

*count \* sizeof(datatype)* bytes of data

# MPI basics

## First six functions (C bindings)

**MPI\_Send (buf, count, datatype, dest, tag, comm)**

*Send message*  
**Assignment Project Exam Help**

*buf* **<https://powcoder.com>** *address of send buffer*

*count* **Add WeChat powcoder** *no. of elements to send ( $\geq 0$ )*

*datatype* *process id of destination*

*dest* *process id of destination*

*tag* *message tag*

*comm* *communicator (handle)*

# MPI basics

## First six functions (C bindings)

**MPI\_Send (buf, count, datatype, dest, tag, comm)**

*Send message*  
**Assignment Project Exam Help**

*buf* **<https://powcoder.com>** *address of send buffer*

*count* **Add WeChat powcoder** *no. of elements to send ( $\geq 0$ )*

*datatype* *process id of destination*

*dest* *process id of destination*

*tag* *message tag*

*comm* *communicator (handle)*

# MPI basics

## First six functions (C bindings)

**MPI\_Send (buf, count, datatype, dest, tag, comm)**

*Send message*  
**Assignment Project Exam Help**

*buf* **<https://powcoder.com>** *address of send buffer*

*count* **Add WeChat powcoder** *no. of elements to send ( $\geq 0$ )*

*datatype* *process id of destination*

*dest* *process id of destination*

*tag* *message tag*

*comm* *communicator (handle)*

# MPI basics

## First six functions (C bindings)

**MPI\_Recv (buf, count, datatype, source, tag, comm, status)**

Assignment Project Exam Help

*Receive a message*

*buf*

*address of receive buffer*  
<https://powcoder.com>

*count*

*max no. of elements in receive buffer ( $\geq 0$ )*

*datatype*

*of receive buffer elements*  
Add WeChat powcoder

*source*

*process id of source process, or  
MPI\_ANY\_SOURCE*

*tag*

*message tag, or MPI\_ANY\_TAG*

*comm*

*communicator*

*status*

*status object*



# MPI basics

## First six functions (C bindings)

### **MPI\_Init (int \*argc, char\*\*argv)**

*Initiate a MPI computation*

*argc(number of arguments) and argv(argument vector) hold main program's arguments*

*Must be called first, and once per process*

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

### **MPI\_Finalize()**

*Shut down a computation*

*The last thing that happens*

# MPI basics

## First six functions (C bindings)

### **MPI\_Comm\_size (MPI\_Comm comm, int\* size)**

*Determine number of processes in comm*  
*comm is communicator handle,*

*MPI\_COMM\_WORLD is the default (including **all** MPI processes)*

***size** holds number of processes in group*

### **MPI\_Comm\_rank (MPI\_Comm comm, int\* pid)**

*Determine id of current (or calling) process*

***pid** holds id of current process*

# MPI basics – a basic example

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[]) {
    int rank, nprocs;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    printf("Hello, world. I am %d of %d\n", rank, nprocs);
    MPI_Finalize();
}
```

mpirun -np 4 myprog

Hello, world. I am 1 of 4

Hello, world. I am 3 of 4

Hello, world. I am 0 of 4

Hello, world. I am 2 of 4

# MPI basics – send and recv example (1)

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[]) {
    int rank, size, i;
    int buffer[10];
    MPI_Status status;
```

Assignment Project Exam Help

```
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (size < 2) {
        printf("Please run with two processes.\n");
        MPI_Finalize();
        return 0;
    }
    if (rank == 0) {
        for (i=0; i<10; i++)
            buffer[i] = i;
        MPI_Send(buffer, 10, MPI_INT, 1, 123, MPI_COMM_WORLD);
    }
```

## MPI basics – send and recv example (2)

```
if (rank == 1) {  
    for (i=0; i<10; i++)  
        buffer[i] = -1;  
    MPI_Recv(buffer, 10, MPI_INT, 0, 123, MPI_COMM_WORLD,  
&status);  
    for (i=0; i<10; i++) {  
        if (buffer[i] != i)  
            printf("Error: buffer[%d] = %d but is expected to be %d\n", i, buffer[i], i);  
    }  
}  
MPI_Finalize();  
}
```



# MPI language bindings

- ❑ Standard (accepted) bindings for Fortran, C and C++
- ❑ Two types of Java bindings (work in progress)
  - ❑ native MPI bindings: native MPI library is called by Java programs through Java wrappers
    - ❑ JavaMPI
    - ❑ mpiJava
  - ❑ pure Java implementations: the whole MPI library is rewritten in Java
    - ❑ jmpj
    - ❑ MPIJ
- ❑ Java Grande Forum trying to sort it all out
- ❑ We will use the C bindings

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder