

High Performance Computing *Course Notes*

Assignment Project Exam Help

<https://powcoder.com>

Shared Memory Parallel
Add WeChat powcoder Programming

Dr Ligang He



OpenMP

- ❑ OpenMP stands for Open specification for Multi-processing
- ❑ used to assist compilers to understand and parallelise the serial code
- ❑ Can be used to specify shared memory parallelism in Fortran, C and C++ programs
- ❑ OpenMP is a specification for
 - ❑ a set of compiler directives,
 - ❑ RUN TIME library routines, and
 - ❑ environment variables

History of OpenMP

- ❑ Started late 80s there was emergence of shared memory parallel computers with proprietary directive-driven programming environments
- ❑ Poor portability, OpenMP emerges as an industry standard
- ❑ OpenMP specifications include:
 - ❑ OpenMP 1.0 for Fortran, 1997, OpenMP 1.0 for C/C++, 1998
 - ❑ OpenMP 2.0 for Fortran, 2000, OpenMP 2.0 for C/C++ , 2002
 - ❑ OpenMP 2.5 for C/C++ and Fortran, 2005
 - ❑ OpenMP 3.0 for C/C++ and Fortran, 2008
 - ❑ OpenMP 3.1, 2011
 - ❑ OpenMP 4.0, 2013, OpenMP 4.5, 2015
- ❑ OpenMP Architecture Review Board: Compaq, HP, IBM, Intel, SGI, SUN

OpenMP programming model

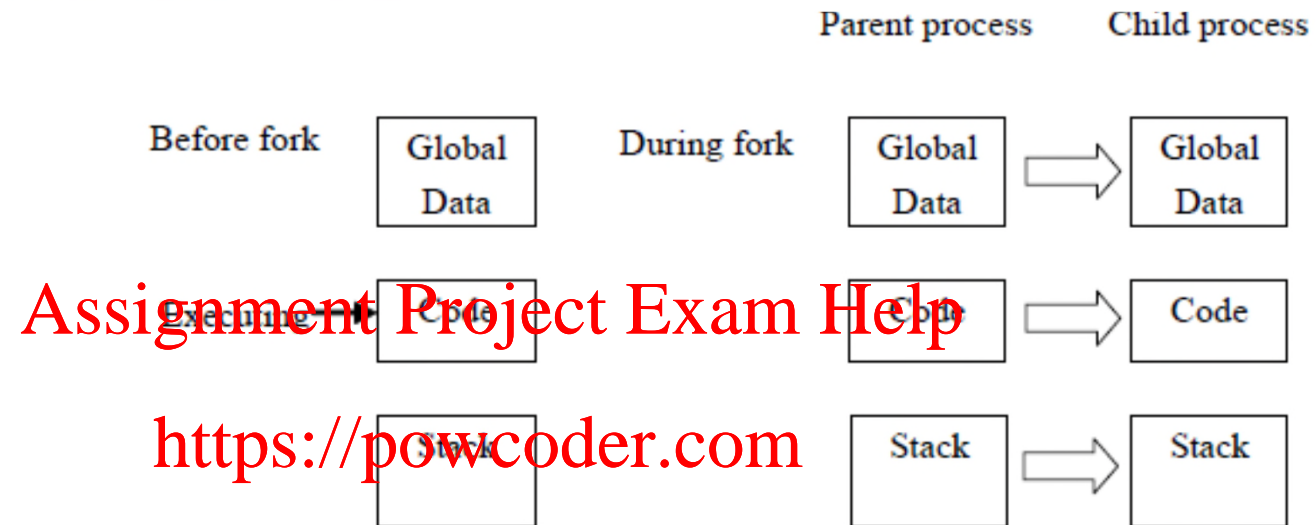
- ❑ An implementation of thread models
 - ❑ Multiple threads running in parallel
- ❑ Used for shared memory architecture
- ❑ Fork-join model

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

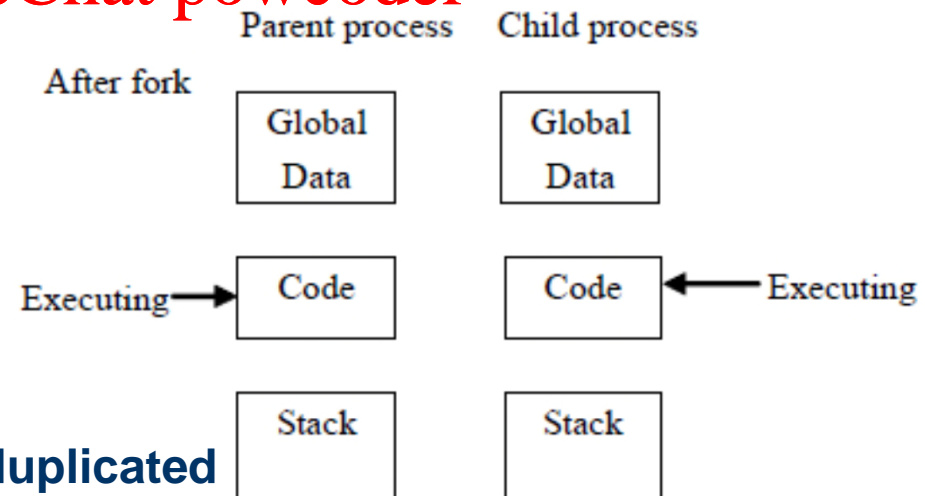
How a new process is created



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



Use the fork function

All three segments and
the program counter are duplicated

How Threads are Created?



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Only the stack segment and the program counter are duplicated

Threads

- ❑ Used to split a program into separate tasks, one per thread, that can execute concurrently
- ❑ “Light weight process”: multiple threads exist within the context of a single process, sharing the process’s code, global information, other resources
- ❑ Threads usually communicate by processing shared global data values

global shared space – global data accessed from single global address space (heap) shared among the threads

local private space – each thread also has its own local private data (stack) that is not shared

OpenMP code structure in C

```
#include <omp.h>
main () {
    int var1, var2, var3;
    Serial code
    /*Beginning of parallel section. Fork a team of threads. Specify variable scoping*/
    #pragma omp parallel private(var1, var2) shared(var3)
    {
        Parallel section: execute by threads
        ...
        All threads join master thread and disband
    }
    Resume serial code
}
```


OpenMP code structure in Fortran

```
PROGRAM HELLO
```

```
INTEGER VAR1, VAR2, VAR3
```

```
Serial code ...
```

```
!Beginning of parallel section. Fork a team of threads. Specify  
variable scoping
```

```
!$OMP PARALLEL PRIVATE(VAR1, VAR2) SHARED(VAR3)
```

```
Parallel section executed by all threads
```

```
...
```

```
All threads join master thread and disband
```

```
!$OMP END PARALLEL
```

```
Resume serial code
```

```
...
```

```
END
```

OpenMP Directives Format

C/C++

<code>#pragma omp</code>	directive-name	[clause, ...]
Required for all OpenMP directives.	A valid OpenMP directive. Must appear after the <code>#pragma omp</code> and before any clauses.	Optional. Clauses can be in any order, and repeated as necessary unless otherwise restricted.

Assignment Project Exam Help
<https://powcoder.com>

Add WeChat powcoder

Fortran

sentinel	directive-name	[clause ...]
All Fortran OpenMP directives must begin with a sentinel. The accepted sentinels depend upon the type of Fortran source. Possible sentinels are: <code>!\$OMP</code> , <code>C\$OMP</code> , <code>*\$OMP</code> .	A valid OpenMP directive. Must appear after the sentinel and before any clauses.	Optional. Clauses can be in any order, and repeated as necessary unless otherwise restricted.

OpenMP features

- ☐ OpenMP directives are ignored by compilers that don't support OpenMP. In this case, codes are run as serial codes
- ☐ Compiler directives used to specify
 - ☐ sections of code that can be executed in parallel
 - ☐ critical sections
 - ☐ Scope of variables (private or shared)
- ☐ Mainly used to parallelize loops, e.g. separate threads to handle separate iterations of the loop
- ☐ There is also a run-time library that has several useful routines for checking the number of threads and number of processors, changing the number of threads, etc

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Fork-Join Model

Multiple threads are created using the parallel construct

For C and C++

```
#pragma omp parallel
```

```
{
```

Assignment Project Exam Help

```
... do stuff
```

```
}
```

<https://powcoder.com>


For Fortran

```
$OMP PARALLEL
```

Add WeChat powcoder

```
... do stuff
```

```
!$OMP END PARALLEL
```



How many threads are generated

The number of threads in a parallel region is determined by the following factors, in order of precedence:

- ☐ Use of the `omp_set_num_threads()` library function
- ☐ Setting of the `OMP_NUM_THREADS` environment variable
- ☐ Implementation default - the number of CPUs on a node

Threads are numbered from 0 (master thread) to N-1

Parallelizing loops in OpenMP

- **Compiler directive specifies that loop can be done in parallel**

For C and C++:

```
#pragma omp parallel for  
for (i=0; i++; i<N) {  
    value[i] = compute(i);  
}
```

For Fortran: **Add WeChat powcoder**

```
!$OMP PARALLEL DO  
  DO (i=1:N)  
    value(i) = compute(i);  
  END DO  
!$OMP END PARALLEL DO
```

Partition and allocation of loop iterations

- Can use thread scheduling to specify partition and allocation of iterations to threads

`#pragma omp parallel for schedule(static,4)`

→ `schedule(static [,chunk])`

→ Partition the loop into blocks of iterations of size chunk

→ Before execution, deal out the blocks to each thread

Partition and allocation of loop iterations

- ❑ Can use thread scheduling to specify partition and allocation of iterations to threads

`#pragma omp parallel for schedule(static,4)`

→ `schedule(static [,chunk])`

<https://powcoder.com>

→ `schedule(dynamic [,chunk])`

→ Partition the loop into blocks of iterations of size chunk

→ Put the blocks into a queue

→ During the execution, Each thread grabs a block off the queue until all are done

→ Which thread runs which block of iterations depends on the thread's execution pace

Partition and allocation of loop iterations

- Can use thread scheduling to specify partition and allocation of iterations to threads

`#pragma omp parallel for schedule(static,4)`

Assignment Project Exam Help

→ `schedule(static [,chunk])`

<https://powcoder.com>

→ `schedule(dynamic [,chunk])`

Add WeChat powcoder

→ `schedule(runtime)`: Find schedule from an environment variable `OMP_SCHEDULE`

Synchronisation in OpenMP

Critical construct

Fortran	<pre>!\$OMP CRITICAL [name] block !\$OMP END CRITICAL</pre>
C/C++	<pre>#pragma omp critical [name] structured_block</pre>

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Barrier construct

Fortran	<pre>!\$OMP BARRIER</pre>
C/C++	<pre>#pragma omp barrier</pre>

Example of Critical Section

```
#include <omp.h>
```

```
main() {
```

```
    int x;
```

```
    x = 0;
```

```
    #pragma omp parallel shared(x)
```

```
    {
```

```
        ...
```

```
        #pragma omp critical
```

```
        x = x+1;
```

```
    } /* end of parallel section */
```

```
}
```

Example of Barrier in OpenMP

```
#include <omp.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    int th_id, nthreads;
    #pragma omp parallel private(th_id)
    {
        th_id = omp_get_thread_num();
        printf("Hello World from thread %d\n", th_id);
        #pragma omp barrier
        if ( th_id == 0 ) {
            nthreads = omp_get_num_threads();
            printf("There are %d threads\n",nthreads);
        }
    }
    return 0;
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Data Scope Attributes in OpenMP

- OpenMP Data Scope Attribute Clauses are used to explicitly define how variables should be viewed by threads
- These clauses are used in conjunction with several directives (e.g. PARALLEL, DO/for) to control the scoping of enclosed variables
- Three often encountered clauses
 - ☐ Shared
 - ☐ Private
 - ☐ Reduction

Shared and private data in OpenMP

- ❑ **private(var)** creates a local copy of var for each thread
- ❑ **shared(var)** states that var is a global variable to be shared among threads

<https://powcoder.com>
Add WeChat powcoder

```
!$OMP PARALLEL DO
!$OMP& PRIVATE(xx,yy) SHARED(u,f)
  DO j = 1,m
    DO i = 1,n
      xx = -1.0 + dx * (i-1)
      yy = -1.0 + dy * (j-1)
      u(i,j) = 0.0
      f(i,j) = -alpha * (1.0-xx*xx) * & (1.0-yy*yy)
    END DO
  END DO
!$OMP END PARALLEL DO
```

Reduction Clause

- ❑ Reduction – `reduction (op : var)`
- ❑ The exemplar op is add, logical OR (commutative operations)
- ❑ A local copy of the variable is made for each thread
- ❑ The local values of the variable can be updated by the threads.
- ❑ At the end of parallel region, the local values are combined to create global value through Reduction operation

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

An Example of Reduction Clause

```
double ZZ, res=0.0;
#pragma omp parallel for reduction (+:res) private(ZZ)
for (i=1;i<=N;i++) {
    ZZ = i;
    res = res + ZZ;
}
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Run-Time Library Routines

→ Can perform a variety of functions, including

Query the number of threads/thread no.

Assignment Project Exam Help

Set the number of threads to be generated

<https://powcoder.com>

Query the number of processors in the computer

Add WeChat powcoder

Changing the number of threads

Run-Time Library Routines

- ❑ query routines allow you to get the number of threads and the ID of a specific thread

`id = omp_get_thread_num(); //thread no.`

`Nthreads = omp_get_num_threads(); //number of threads`

- ❑ Can specify number of threads at runtime

`omp_set_num_threads(Nthreads);`

Environment Variable

→ Controlling the execution of parallel code

→ Four environment variables

- Assignment Project Exam Help**
<https://powcoder.com>
- ☐ **OMP_SCHEDULE:** how iterations of a loop are scheduled
 - ☐ **OMP_NUM_THREADS:** maximum number of threads
 - ☐ **OMP_DYNAMIC:** enable or disable dynamic adjustment of the number of threads
 - ☐ **OMP_NESTED:** enable or disable nested parallelism

Lab session today

- Practice OpenMP
- Download lab instructions and code from here:

https://warwick.ac.uk/fac/sci/dcs/teaching/material/cs402/cs402_seminar2_openmp.pdf

https://warwick.ac.uk/fac/sci/dcs/teaching/material/cs402/cs402_seminar2_code.zip

- Move down to Lab 001 and 003

<https://powcoder.com>
Add WeChat powcoder

Assignment 1 - OpenMP

- Use OpenMP to parallelize the deqn code
 - The overall objective is to achieve good speedup
- Write a report
 - Explain in detail what you did with the sequential code
 - benchmark the runtime of each relevant loop and the runtime of the whole parallel program against the number of threads; present the runtimes in graph or table; analyze the results
 - Discuss the iteration scheduling in your program
 - Analyze the overhead of OpenMP
 - Presentation skills, spelling, punctuation and grammar
 - Up to four A4 pages

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder